

Concurrency

Review: Concurrency

- Concurrent access to shared data may lead to synchronization errors
- For example:

```
void increment(int* ip)
{
    *ip = *ip + 1;
}
```

Concurrency

- The C code from the previous slide compiles into the following assembly:

```
increment:
    pushl %ebx
    movl 8(%esp), %eax # %eax = ip
    movl (%eax), %ebx  # %ebx = *ip
    add 1, %ebx        # %ebx = *ip + 1
    movl %ebx, (%eax)  # *ip = *ip + 1
    popl %ebx
    ret
```

Concurrency

- So far TOS only implements cooperative multi-tasking. I.e., a process must relinquish control voluntarily via a call to `resign()`
- In pre-emptive multitasking, a context switch can occur between any two machine instructions!
- In this case, concurrency issues may lead to what is called *race conditions*.
- TOS will soon support pre-emptive multi-tasking, so we have to learn how to deal with such problems.

Concurrency

- For the following two scenarios, we assume global variables as follows:

```
int x = 5;
```

```
int* p = &x;
```

```
int* q = &x;
```

- Let address of `x` be `0x5000`
- Process 1 executes: `increment(p)`
- Process 2 executes: `increment(q)`

Concurrency – Scenario 1

Process 1

Process 2

Time
↓

```
movl 8(%esp), %eax  
movl (%eax), %ebx
```

```
movl 8(%esp), %eax
```

%eax: 0x5000

%ebx: 5

%eax: 0x5000

%ebx: XXXXXXXX

Memory address 0x5000: 5

Concurrency – Scenario 1

Process 1

Process 2

```
movl 8(%esp), %eax
movl (%eax), %ebx
add 1, %ebx
```

```
movl 8(%esp), %eax
```

%eax: 0x5000

%ebx: 6

%eax: 0x5000

%ebx: XXXXXXXX

Memory address 0x5000: 5

Concurrency – Scenario 1

Process 1

Process 2

```
movl 8(%esp), %eax
movl (%eax), %ebx
add 1, %ebx
movl %ebx, (%eax)
```

```
movl 8(%esp), %eax
```

%eax: 0x5000

%ebx: 6

%eax: 0x5000

%ebx: XXXXXXXX

Memory address 0x5000: 6

Concurrency – Scenario 1

Process 1

Process 2

```
movl 8(%esp), %eax
movl (%eax), %ebx
add 1, %ebx
movl %ebx, (%eax)
```

```
movl 8(%esp), %eax
```

```
movl (%eax), %ebx
```

%eax: 0x5000

%ebx: 6

%eax: 0x5000

%ebx: 6

Memory address 0x5000: 6

Concurrency – Scenario 1

Process 1

```
movl 8(%esp), %eax
movl (%eax), %ebx
add 1, %ebx
movl %ebx, (%eax)
```

%eax: 0x5000

%ebx: 6

Process 2

```
movl 8(%esp), %eax
```

```
movl (%eax), %ebx
```

```
add 1, %ebx
```

```
movl %ebx, (%eax)
```

%eax: 0x5000

%ebx: 7

Memory address 0x5000: 7

Concurrency – Scenario 2

Process 1

Process 2

```
movl 8(%esp), %eax  
movl (%eax), %ebx
```

```
movl 8(%esp), %eax
```

%eax: 0x5000
%ebx: 5

%eax: 0x5000
%ebx: XXXXXXXX

Memory address 0x5000: 5

Concurrency – Scenario 2

Process 1

```
movl 8(%esp), %eax  
movl (%eax), %ebx
```

Process 2

```
movl 8(%esp), %eax  
  
movl (%eax), %ebx
```

%eax: 0x5000
%ebx: 5

%eax: 0x5000
%ebx: 5

Memory address 0x5000: 5

Concurrency – Scenario 2

Process 1

```
movl 8(%esp), %eax  
movl (%eax), %ebx
```

Process 2

```
movl 8(%esp), %eax
```

```
movl (%eax), %ebx  
add 1, %ebx  
movl %ebx, (%eax)
```

%eax: 0x5000

%ebx: 5

%eax: 0x5000

%ebx: 6

Memory address 0x5000: 6

Concurrency – Scenario 2

Process 1

```
movl 8(%esp), %eax  
movl (%eax), %ebx
```

```
add 1, %ebx
```

Process 2

```
movl 8(%esp), %eax
```

```
movl (%eax), %ebx  
add 1, %ebx  
movl %ebx, (%eax)
```

%eax: 0x5000

%ebx: 6

%eax: 0x5000

%ebx: 6

Memory address 0x5000: 6

Concurrency – Scenario 2

Process 1

```
movl 8(%esp), %eax
movl (%eax), %ebx
```

```
add 1, %ebx
movl %ebx, (%eax)
```

```
%eax: 0x5000
%ebx: 6
```

Process 2

```
movl 8(%esp), %eax
```

```
movl (%eax), %ebx
add 1, %ebx
movl %ebx, (%eax)
```

```
%eax: 0x5000
%ebx: 6
```

Memory address 0x5000: 6

Race Conditions

- Scenario 1 executes as expected.
- Scenario 2 leads to a so-called *race condition* because context switches happen at unfortunate moments.
- It is called race condition, because of a “race” between two processes.
- Race conditions only occur rarely, but are very difficult to debug.
- A pre-condition for a race condition is that two processes must access a shared resource (e.g., the same global variable).

Fixing Concurrency Bugs

- To fix this problem, we can use a *lock*.
- Operations on a lock: acquire and release
- When one task acquires a lock, no other task may acquire it until the first task calls release.
 - In other words, only one task at a time may hold the lock

Train Semaphore



- Semaphore signals train if it is safe to enter a “critical section”.

Fixing Concurrency Bugs

```
void increment(int* ip, lock* l)
{
    acquire(l);
    *ip = *ip + 1;
    release(l);
}
```

- Now, a task that begins to increment `*ip` must finish before another task may begin

Implementing Locks

- Need help from the hardware
- Instructions are *atomic*: once an instruction begins executing, nothing else happens until it is finished.

Implementing Locks

- Every modern architecture provides some useful primitives for implementing locks.
- Atomic test-and-set:
 - Test a value (e.g., is value == 0) and set it in a single atomic operation
- Intel x86 also provides atomic swap and atomic load-compute-store (xchg).
- **Conceptually, `xchg %eax, (memaddr)` does the following:**

```
pushl %ebx           # save %ebx
movl  (memaddr), %ebx # %ebx = *memaddr
pushl %ebx           # swap %eax and %ebx
pushl %eax
popl  %ebx
popl  %eax
movl  %ebx, (memaddr) # *memaddr = %ebx
popl  %ebx           # restore %ebx
```

Spin Locks

```
lock:    dd 0                # The lock variable. 1 = locked, 0 = unlocked.
```

```
spin_acquire:
```

```
    mov $1, %eax            # Set the EAX register to 1.
loop:    xchg %eax, (lock)    # Atomically swap the EAX register with
                             # the lock variable. This will always
                             # store 1 to the lock, leaving previous
                             # value in the EAX register.
    test %eax, %eax         # Test EAX with itself. Among other
                             # things, this will set the processor's
                             # Zero Flag if EAX is 0. If EAX is 0,
                             # then the lock was unlocked and we just
                             # locked it. Otherwise, EAX is 1 and we
                             # didn't acquire the lock.
    jnz loop                # Jump back to the XCHG instruction if
                             # the Zero Flag is not set, the lock was
                             # locked, and we need to spin.
    ret                     # The lock has been acquired, return to
                             # the calling function.
```

```
spin_release:
```

```
    mov $0, %eax            # Set the EAX register to 0.
    xchg %eax, (lock)        # Atomically swap the EAX register with
                             # the lock variable.
    ret                     # The lock has been released.
```

Spin Locks

- On the previous slide, the code tests a memory location (`lock`). If this memory location contains a 1, it means another process has already obtained the lock. If the memory location is 0, it means the lock is available. The atomic `xchg` instruction is used to attempt to do an exchange of 1 with the memory location. If `%eax` contains 0 after the `xchg` instruction, it means that the lock was achieved by the current process. If the `%eax` contains a 1 after the atomic `xchg` instruction this signifies that another process already has the lock.

Building a Better Lock

- The problem with spin locks: during a lengthy critical section, other tasks waste CPU cycles (which is called *busy wait*).
- Spin locks are great for short critical sections.
- For longer critical sections, we want a lock that will cause the task to go to “sleep” if the lock is not available.
- “Sleep”: process is off the ready queue.

Building a Better Lock

```
struct lock {  
    enum { HELD, AVAILABLE } status;  
    PROCESS waiting;  
}  
  
void acquire(struct lock* l)  
{  
    if (l->status != AVAILABLE) {  
        append(l->waiting, active_proc);  
        remove_ready_queue(active_proc);  
    }  
    l->status = HELD;  
}
```

Locks

- Problem: race condition inside `acquire()`
 - If there is a context switch after we check `status` but before changing it, two processes hold the lock simultaneously!
- Solution: the critical section inside `acquire()` is short, so use a spin lock

```
struct lock {  
    spinlock slock;  
    enum { HELD, AVAILABLE } status;  
    PROCESS waiting;  
}
```

Locks - Acquire

```
void acquire(struct lock* l)
{
    spin_acquire(l->slock);
    while (l->status != AVAILABLE) {
        append(l->waiting, active_proc);
        spin_release(l->slock);
        remove_ready_queue(active_proc);
        spin_acquire(l->slock);
    }
    l->status = HELD;
    spin_release(l->slock);
}
```

Locks - Release

```
void release(struct lock* l)
{
    spin_acquire(l->slock);
    l->status = AVAILABLE;
    spin_release(l->slock);
    add_ready_queue(l->waiting);
}
```