

Segmentation

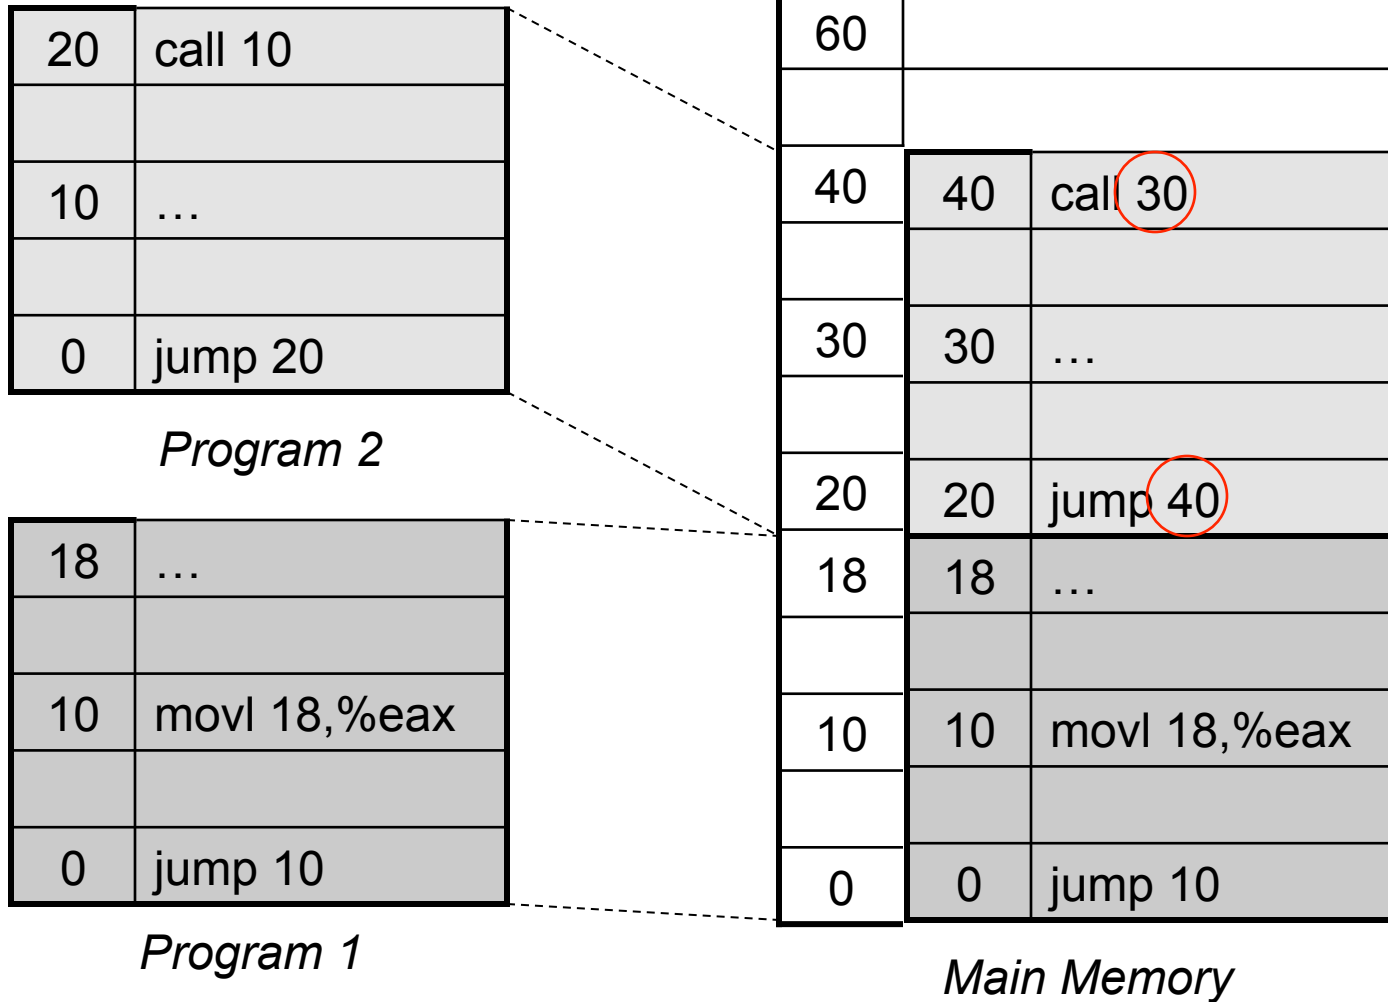
Objective

- Explain the x86 segmentation model
- Explain how a virtual address is translated by the x86 to a physical address
- Explain the various x86 datastructures and hardware registers

Loading of Programs (1)

- Every program is linked by the linker to be a self-contained executable.
- Programs are linked as if they were loaded to address 0. I.e., all references (jumps, calls, memory loads) are done relative to address 0.
- When a program is loaded into memory, it will be assigned an available region of memory.
- All references need to be adjusted to the base address of this memory region (see next slide).
- Beware: some references (such as 0xB8000) should not be adjusted!
- Note: adjusting references does not guard against one process corrupting another process' memory!

Loading of Programs (2)



Address Translation

- “Patching” a program as shown on the previous slide is a complicated matter and also does not protect processes amongst each other.
- A better way would be to have hardware support that does *not* require this patching.
- This is the goal of *address translation*: instead of patching a program manually, (virtual) addresses are translated ‘on the fly’.
- Physical memory in a computer is a linear array of bytes called *physical address space*.
- Physical memory is not directly accessed by programs. They use virtual memory which maps to physical memory.
- Chief benefit of virtual memory is protection as each program (process) gets its own *address space*. A program can not access address space of another program.

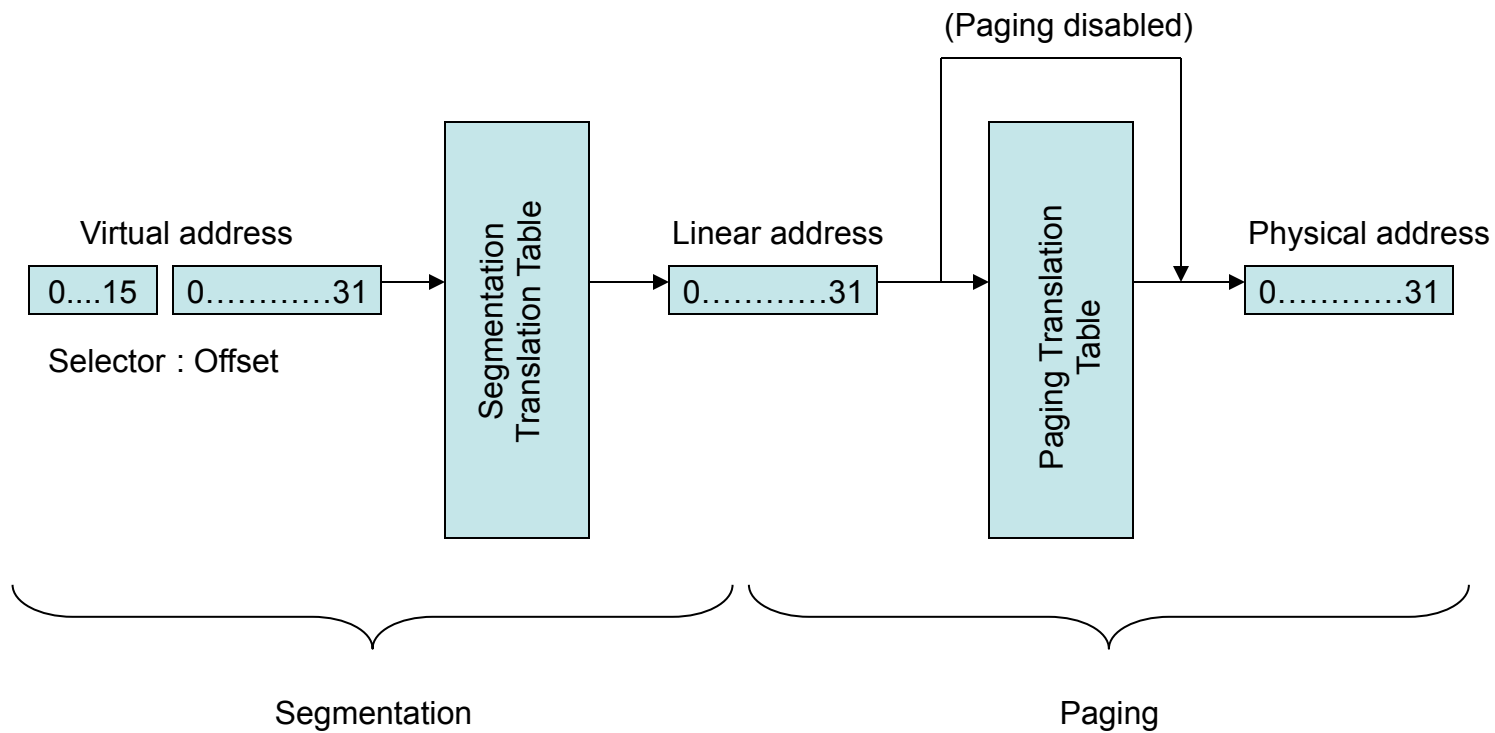
Segmentation and Paging - I

- Segmentation and paging are two *address translation techniques* supported by the x86.
- Segmentation is a “two dimensional” *virtual address space* defined by selector:offset.
- Each selector value points to a segment and the offset value points to the offset within that segment
 - Selector: 16 bits
 - Offset: 32 bits
- Segmentation maps *virtual addresses* to *linear addresses*.
- All x86 programs have to go through segmentation.

Segmentation and Paging - II

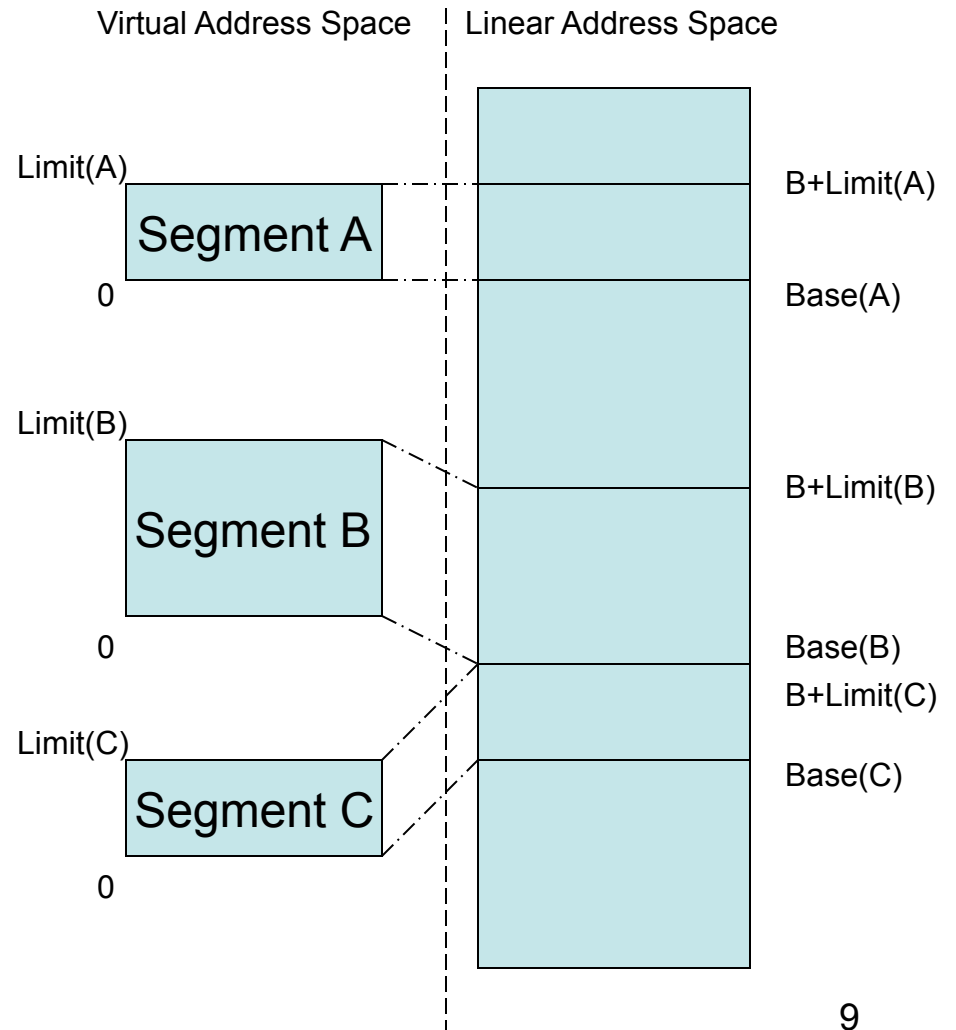
- Linear address space is 32 bit address space for each task
- Allows up to a maximum of 4 GB ($= 2^{32}$ bytes)
- Paging maps linear address space to physical address space
- Paging can be enabled or disabled (we will not use paging!)
- If paging is disabled, then the linear address corresponds to the physical address
- Both segmentation and paging use translation tables

Segmentation and Paging - III



Segmentation - Mapping

- Variable size units of memory called segments form the basis of the virtual-to-linear address translation
- Segments are defined by:
 - base address
 - address limit
 - segment attributes
- Addresses within one segment are relative to the base address of that segment.
- Idea: load each program into its own segment!



Segmentation - Example

- Segment 1 is defined as follows:
 - Base address: 0x50000
 - Limit: 0xffff

Virtual address	Linear address
1:0	0x50000
1:0x67	0x50067
1:0x10000	Illegal. Beyond limit of segment

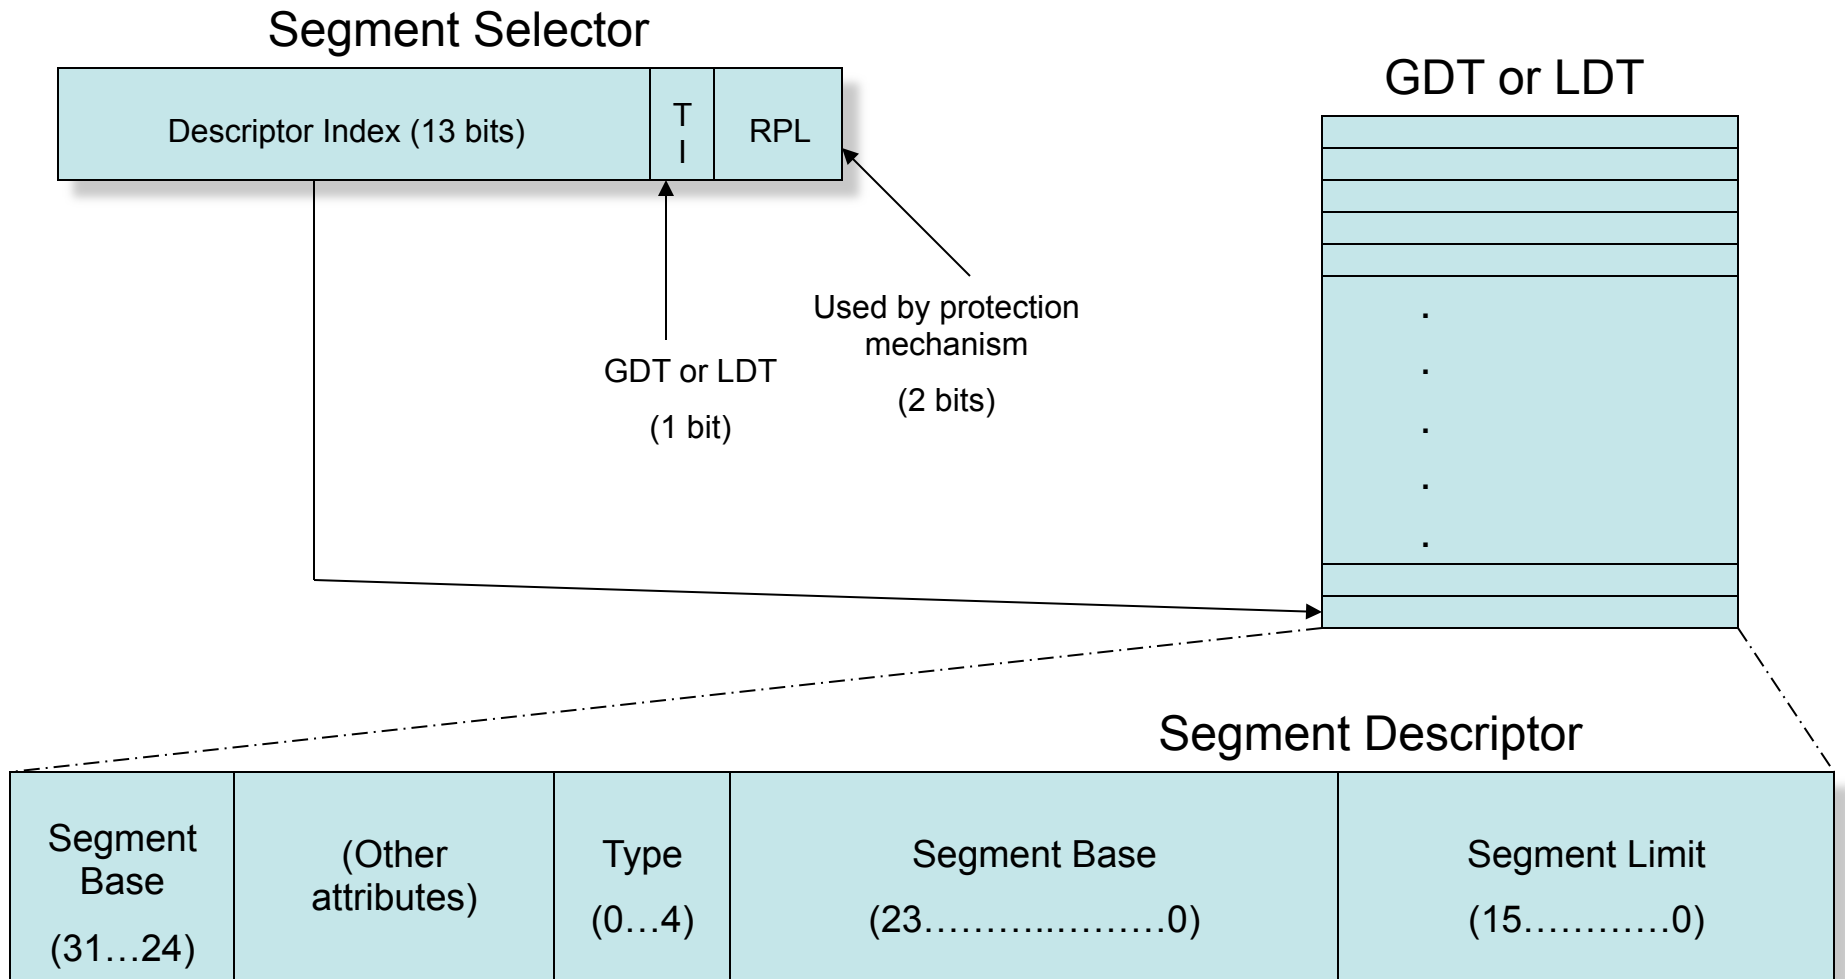
Segment Descriptors

- The x86 contains six *segment registers*: %CS, %DS, %ES, %FS, %GS, and %SS
- A segment register contains a 16 bit segment selector
- Segment selector is an index to the *segment descriptor* stored in a *descriptor table*
- Each segment descriptor contains important information about the segment: base, limit, attributes
- The 4 bit type attribute specifies if the segment is read only/read write/execute or other specific combinations thereof

Segment Descriptor Tables

- Segment descriptor tables hold the segment descriptors
- There are two types local (LDT) and global (GDT). The virtual address space is distributed between these two tables
- GDT remains the same for all processes, while the LDT is private to each process
- GDT can be used to store information about OS code and data segments accessible to all programs. LDT can be used to store code and data segments for that program
- GDT also stores other types of descriptors (not relevant to the topic here)
- In TOS we only use the GDT

Anatomy of a Segment Selector and Descriptor




Notes

- Size of one Segment Descriptor is 8 bytes
- Note that the 32 bits of the segment base are not stored consecutively in the Segment Descriptor
- RPL (Requested Privilege Level). Should always be set to 0 in TOS
- TI (Table Index). Should always be set to 0 in TOS
- The only two segment types we will use in TOS are:
 - Data: 0x2
 - Code: 0xa
- GDT and LDT are tables which are stored in ordinary memory

LGDT instruction

- Load GDT instruction (lgdt) loads the 48 bit GDTR register with the base address and limit of the GDT.
- This instruction is used in TOS' boot loader



Segmentation and X86 Assembly

- How are segments used in day-to-day assembly?
- Each kind of memory access uses a certain default segment register
 - `%CS` is used for fetching from memory (e.g. `CALL`, `JMP`, `RET`)
 - `%SS` is used for all memory access to the stack (e.g. `PUSH`, `POP`)
 - `%DS` is used for all other memory access (e.g. `MOV`)
- Examples:
 - `PUSH $1` access `%SS:%ESP`
 - `MOV $1, (%EAX)` access `%DS:(%EAX)`
 - `JMP $100` access `%CS:$100`
- The content of segment registers are typically loaded only once at boot-time and then never changed again
- Segment registers can be loaded from normal x86 register, e.g.

```
movw    $0x10, %AX
movw    %AX, %DS
```

loads the `%DS` segment register with 16

GDT for TOS (1)

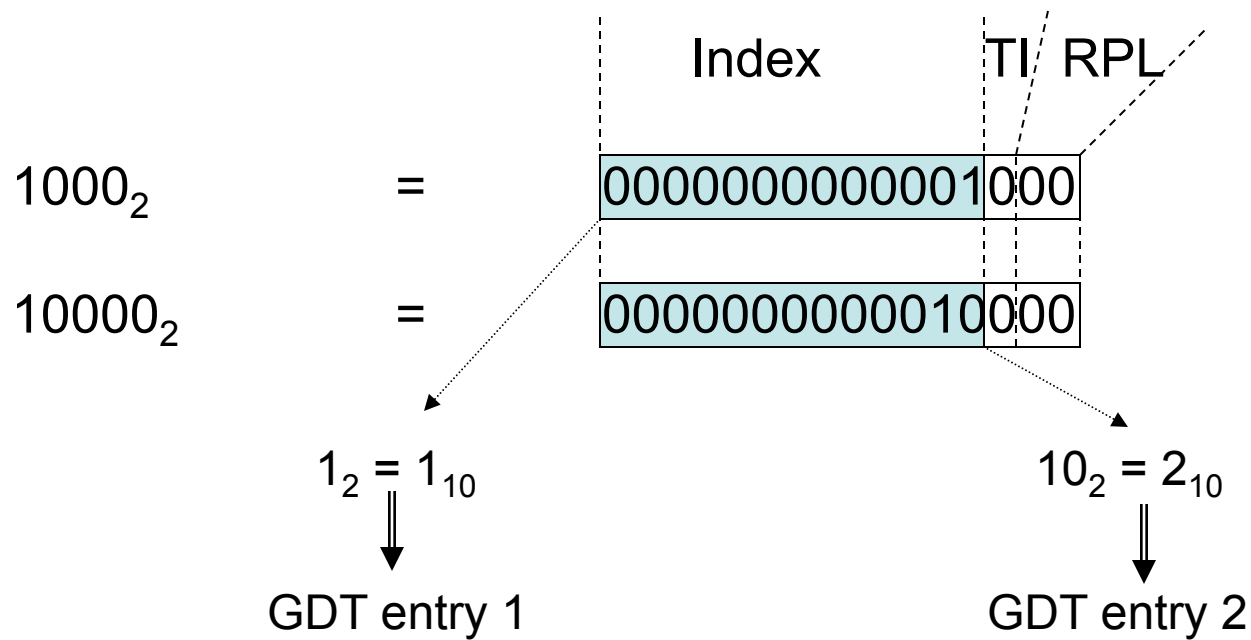
- Basic idea: segments in TOS are defined such that a virtual address is identical to the physical address.
- The GDT for TOS is constructed and loaded during the boot process; i.e., before calling `kernel_main()`
- Size of the GDT is 24 bytes ($= 3 * 8$ bytes).
- `%CS` is loaded with 0x8 (GDT entry 1) and `%DS` and `%SS` are loaded with 0x10 (GDT entry 2)

Entry	Base	Limit	Attributes	Comments
0	-	-	-	Dummy Entry
1	0	0xFFFFFFFF	CODE	Used for <code>%CS</code>
2	0	0xFFFFFFFF	DATA	Used for <code>%DS</code> and <code>%SS</code>

GDT for TOS (2)

- TOS uses two segment selectors:

CODE: $8_{16} = 1000_2$
DATA: $10_{16} = 10000_2$



Memory Access in TOS

- Given the following code:

```
movl    $0xB8000, %EAX
movb    $'A', (%EAX)
```

- 'A' is stored to 0xB8000. But how does it look exactly with segmentation?
- Since we do a regular memory access, the %DS segment is used. The linear address is therefore 0x10: 0xB8000
 - Segment 0x10
 - Offset 0xB8000
- Base address of the segment 0x10 is 0
- Therefore the linear address is $0xB8000 + 0 = 0xB8000$
- Since we don't use paging, linear address corresponds to physical address
- Therefore, we access physical address 0xB8000

Inter-Segment Subroutines (1)

- So far, when doing a `CALL` instruction, we only specified the 32-bit offset, but not the segment sector
- The segment was implicitly selected through the `%CS` segment register
- This is called an *Intra-Segment Jump* because the jump happens within the same segment
- An *Inter-Segment Jump* jumps between different segments
- Inter-Segment Jumps will happen in TOS only for interrupts
- For an *Inter-Segment Subroutine* call, not only the return address is pushed on the stack, but also `%CS` (see next slide)

Inter-Segment Subroutines (2)

- Assumptions:
 - %CS = 0x1234
 - Return address is 0xABCD1234
- When executing `CALL 0x5678: 0x12123434` the following information is pushed onto the stack:
 - 0x00001234 (old value of %CS as a 32-bit value)
 - 0xABCD1234 (return address)
- For an intra-segment subroutine call only the offset part 0xABCD1234 would have been pushed on the stack
- After pushing the return address on the stack, the registers are loaded as follows:
 - %CS := 0x5678 (this goes through the GDT!)
 - %EIP := 0x12123434